

Breast_Cancer_Diagnosis

November 15, 2022

1 Overview

This is my solution from the course project of IBM [Supervised Machine Learning: Classification](#) course. The Dataset used was the [Breast Cancer Wisconsin \(Diagnostic\)](#) downloaded from Kaggle.

The dataset contains characteristics of cell nuclei obtained from digitized images of breast masses which are: - radius (mean of distances from center to points on the perimeter - texture (standard deviation of gray-scale values) - perimeter - area - smoothness (local variation in radius lengths) - compactness ($\text{perimeter}^2 / \text{area} - 1.0$) - concavity (severity of concave portions of the contour) - concave points (number of concave portions of the contour) - symmetry - fractal dimension ("coastline approximation" - 1)

For each feature was computed the mean, standard error, and "worst" or largest (mean of the three largest values). In addition, the dataset contains an ID for each patient and the type of cancer (M = malignant, B = benign).

The main objective is: **train a classification model to predict the cancer diagnosis based on these cell characteristics**. I also interested on interpretability to understand which features contributes to a better diagnosis.

These and other analyzes related to the [IBM Machine Learning Professional Certificate](#) can be accessed in this [repository](#).

2 Necessary packages

```
[ ]: #Core
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os

#Auto EDA -- !pip install dataprep --
#from dataprep.eda import create_report

#Machine Learning
from sklearn.model_selection import train_test_split, StratifiedKFold, \
    cross_val_score
from sklearn.preprocessing import StandardScaler
```

```

from sklearn.linear_model import LogisticRegressionCV, LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.pipeline import make_pipeline
from sklearn.metrics import roc_auc_score, classification_report, \
    confusion_matrix, ConfusionMatrixDisplay
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb

#Options
pd.set_option('display.max_columns', None)
pd.set_option('display.precision', 3)

```

3 Glimpse on Data

```

[ ]: #loading dataframe
PATH = "data/data.csv"
df_raw = pd.read_csv(os.path.join(PATH))

#create a copy to avoid edit raw data
df = df_raw.copy()

```

```

[ ]: df.head()

```

```

[ ]:
      id diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean \
0   842302         M      17.99      10.38      122.80      1001.0
1   842517         M      20.57      17.77      132.90      1326.0
2  84300903         M      19.69      21.25      130.00      1203.0
3  84348301         M      11.42      20.38       77.58       386.1
4  84358402         M      20.29      14.34      135.10      1297.0

      smoothness_mean  compactness_mean  concavity_mean  concave points_mean \
0           0.118           0.278           0.300           0.147
1           0.085           0.079           0.087           0.070
2           0.110           0.160           0.197           0.128
3           0.142           0.284           0.241           0.105
4           0.100           0.133           0.198           0.104

      symmetry_mean  fractal_dimension_mean  radius_se  texture_se  perimeter_se \
0           0.242           0.079           1.095           0.905           8.589
1           0.181           0.057           0.543           0.734           3.398
2           0.207           0.060           0.746           0.787           4.585
3           0.260           0.097           0.496           1.156           3.445
4           0.181           0.059           0.757           0.781           5.438

      area_se  smoothness_se  compactness_se  concavity_se  concave points_se \

```

0	153.40	0.006	0.049	0.054	0.016
1	74.08	0.005	0.013	0.019	0.013
2	94.03	0.006	0.040	0.038	0.021
3	27.23	0.009	0.075	0.057	0.019
4	94.44	0.011	0.025	0.057	0.019

	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	\
0	0.030	0.006	25.38	17.33	
1	0.014	0.004	24.99	23.41	
2	0.022	0.005	23.57	25.53	
3	0.060	0.009	14.91	26.50	
4	0.018	0.005	22.54	16.67	

	perimeter_worst	area_worst	smoothness_worst	compactness_worst	\
0	184.60	2019.0	0.162	0.666	
1	158.80	1956.0	0.124	0.187	
2	152.50	1709.0	0.144	0.424	
3	98.87	567.7	0.210	0.866	
4	152.20	1575.0	0.137	0.205	

	concavity_worst	concave points_worst	symmetry_worst	\
0	0.712	0.265	0.460	
1	0.242	0.186	0.275	
2	0.450	0.243	0.361	
3	0.687	0.258	0.664	
4	0.400	0.163	0.236	

	fractal_dimension_worst	Unnamed: 32
0	0.119	NaN
1	0.089	NaN
2	0.088	NaN
3	0.173	NaN
4	0.077	NaN

```
[ ]: df.tail()
```

```
[ ]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
564	926424	M	21.56	22.39	142.00	1479.0	
565	926682	M	20.13	28.25	131.20	1261.0	
566	926954	M	16.60	28.08	108.30	858.1	
567	927241	M	20.60	29.33	140.10	1265.0	
568	92751	B	7.76	24.54	47.92	181.0	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
564	0.111	0.116	0.244	0.139	
565	0.098	0.103	0.144	0.098	
566	0.085	0.102	0.093	0.053	

567	0.118	0.277	0.351	0.152
568	0.053	0.044	0.000	0.000

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	\
564	0.173	0.056	1.176	1.256	
565	0.175	0.055	0.765	2.463	
566	0.159	0.056	0.456	1.075	
567	0.240	0.070	0.726	1.595	
568	0.159	0.059	0.386	1.428	

	perimeter_se	area_se	smoothness_se	compactness_se	concavity_se	\
564	7.673	158.70	0.010	0.029	0.052	
565	5.203	99.04	0.006	0.024	0.040	
566	3.425	48.55	0.006	0.037	0.047	
567	5.772	86.22	0.007	0.062	0.071	
568	2.548	19.15	0.007	0.005	0.000	

	concave points_se	symmetry_se	fractal_dimension_se	radius_worst	\
564	0.025	0.011	0.004	25.450	
565	0.017	0.019	0.002	23.690	
566	0.016	0.013	0.004	18.980	
567	0.017	0.023	0.006	25.740	
568	0.000	0.027	0.003	9.456	

	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
564	26.40	166.10	2027.0	0.141	
565	38.25	155.00	1731.0	0.117	
566	34.12	126.70	1124.0	0.114	
567	39.42	184.60	1821.0	0.165	
568	30.37	59.16	268.6	0.090	

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	\
564	0.211	0.411	0.222	0.206	
565	0.192	0.322	0.163	0.257	
566	0.309	0.340	0.142	0.222	
567	0.868	0.939	0.265	0.409	
568	0.064	0.000	0.000	0.287	

	fractal_dimension_worst	Unnamed: 32
564	0.071	NaN
565	0.066	NaN
566	0.078	NaN
567	0.124	NaN
568	0.070	NaN

```
[ ]: df.shape
```

```
[ ]: (569, 33)
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                 569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                               569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                   569 non-null    float64
22  radius_worst                          569 non-null    float64
23  texture_worst                         569 non-null    float64
24  perimeter_worst                       569 non-null    float64
25  area_worst                            569 non-null    float64
26  smoothness_worst                      569 non-null    float64
27  compactness_worst                     569 non-null    float64
28  concavity_worst                       569 non-null    float64
29  concave points_worst                  569 non-null    float64
30  symmetry_worst                        569 non-null    float64
31  fractal_dimension_worst               569 non-null    float64
32  Unnamed: 32                           0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

```
[ ]: #checking id
len(df.id.unique()) == df.shape[0]
```

```
[ ]: True
```

```
[ ]: #checking diagnosis
df.diagnosis.unique()
```

```
[ ]: array(['M', 'B'], dtype=object)
```

- Id contains a unique value for each observation. This kind of feature is not useful for Machine Learning models. I delete it.
- Diagnosis is the unique not numerical feature. I encode it with 0 mean “B” and 1 mean “M”.
- Unnamed: 32 has only missing data.

```
[ ]: #preprocessing
df = df.drop(["id", "Unnamed: 32"], axis = 1)
df.diagnosis = df.diagnosis.map({"B":0, "M":1})
```

I always like to start the analysis with an auto EDA tool to give a first look on data and and get quick insights. Here I use [dataprep](#) for this purpose.

```
[ ]: #this code export a html file with dataprep output which can be opened locally
↳by any browser
report = create_report(df, title='Cancer Detection')
report.save("Cancer Detection.html")
```


Report has been saved to Cancer Detection.html!

Here are some discoveries: - The target class is unbalanced (62.74% Benign and 37.26% Malign) - There are many features with skewness. - There are several features (Radius, Perimeter, Area, Compactness, Concavity) that seem to separate the classe well. I could bet that a simple linear classifier like a Logistic Regression would do a good job devide the the separability of classes. - For that features a great value in them are associated to malign cancer. - Many features with high degree of correlation (>70%) with target - There are high degree of multicollinearity between predictor features. Eliminate them could lead to a more simple model and simple to interpret due to the fewer features. - There are some outilers, two observation in concavity worst for example are very far from others


4 Exploratory Data Analysis

4.1 Bivariate Analysis


```
[ ]: #_mean pairplot
sns.pairplot(data = df , x_vars =[i for i in df.columns if "mean" in i],
↳y_vars = ["diagnosis"]);
```



```
[ ]: #_se pairplot
sns.pairplot(data = df , x_vars =[i for i in df.columns if "_se" in i], y_vars=["diagnosis"]);
```



```
[ ]: #_worst pairplot
sns.pairplot(data = df , x_vars =[i for i in df.columns if "_worst" in i], y_vars=["diagnosis"]);
```



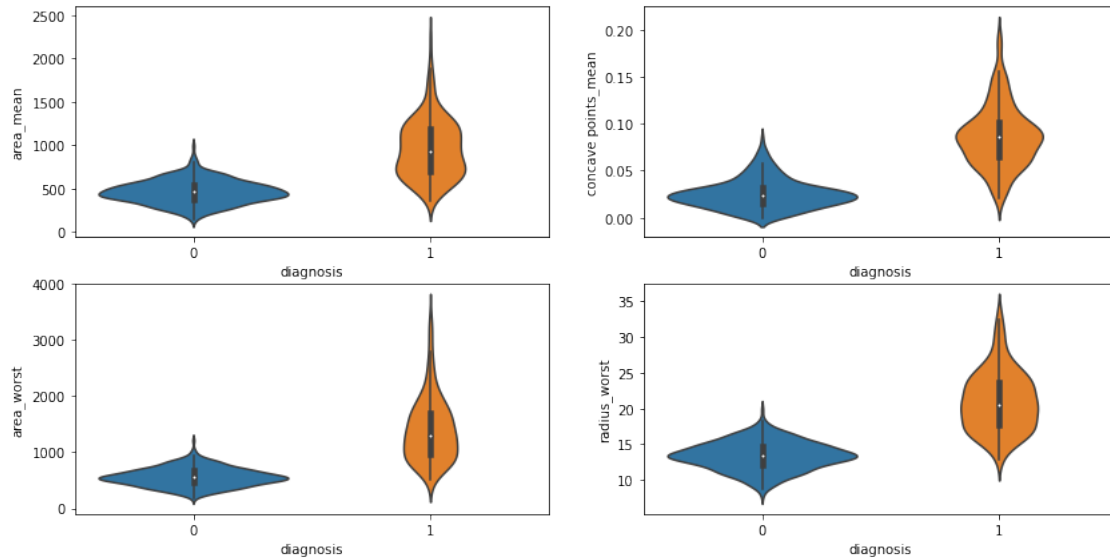
The pairplot confirm the separability between classes usually with large values leading to a malign cancer. “area_mean”, “concave_points_mean”, “area_worst”, “radius_worst” seem to separate the classes very well.

Also, we can see in the _se pairplot that there are two points that are very far from each other.

```
[ ]: to_plot = ["area_mean", "concave points_mean", "area_worst", "radius_worst"]

fig, ax = plt.subplots(2,2, figsize = (14,7))

for i, var in enumerate(to_plot):
    sns.violinplot(data = df, x = "diagnosis", y = var, ax = ax.flat[i])
```



Again, we can see that those features have lower median and mean to Benign and in Benign the concentration of values are in lower values

4.2 Dealing With Outliers on __SE features

Here, I manually looked for the outliers based on pairplot

```
[ ]: df.query("radius_se > 2 and perimeter_se > 15 and area_se > 400").index
```

```
[ ]: Int64Index([212, 461], dtype='int64')
```

```
[ ]: df.query("smoothness_se > 0.02 and diagnosis == 1").index
```

```
[ ]: Int64Index([122, 213], dtype='int64')
```

```
[ ]: df.query("texture_se > 4").index
```

```
[ ]: Int64Index([192], dtype='int64')
```

```
[ ]: df.query("compactness_se > 0.12").index
```

```
[ ]: Int64Index([190], dtype='int64')
```

```
[ ]: df.query("concavity_se > 0.3").index
```

```
[ ]: Int64Index([68, 152], dtype='int64')
```

```
[ ]: df[df["concave points_se"] > 0.05].index
```



```
[ ]: Int64Index([152], dtype='int64')
```

```
[ ]: df.query("(symmetry_se > 0.06 and diagnosis == 0) or symmetry_se > 0.065" ).  
      ↪ index
```

```
[ ]: Int64Index([78, 314], dtype='int64')
```


```
[ ]: df.query("fractal_dimension_se > 0.02").index
```

```
[ ]: Int64Index([71, 152, 290], dtype='int64')
```

```
[ ]: outliers = [212,461,122,192,213,190,68,152,78,314,71,152,290]
```

```
[ ]: df = df.drop(outliers, axis = 0)  
df = df.reset_index(drop = True)
```

```
[ ]: # re-plot the _se pairplot  
sns.pairplot(data = df , x_vars =[i for i in df.columns if "_se" in i], y_vars_  
      ↪ ["diagnosis"]);
```




4.3 Correlation

```
[ ]: corr = abs(df.corr()['diagnosis']).sort_values(ascending = False)[1:]  
corr[corr > 0.75]
```

```
[ ]: concave points_worst    0.797  
perimeter_worst           0.787  
concave points_mean       0.784  
radius_worst              0.779  
Name: diagnosis, dtype: float64
```

```
[ ]: fig, ax = plt.subplots(2,2, figsize = (14,7))  
  
for i, var in enumerate(corr[corr > 0.75].index):  
    sns.violinplot(data = df, x = "diagnosis", y = var, ax = ax.flat[i])  
  
fig.suptitle("Violinplot of high correlated features with diagnosis");
```

Violinplot of high correlated features with diagnosis



5 Create a Baseline

I will create a simple baseline with logistic regression regularized (Lasso) to deal with multicollinearity.

As lasso regularization tends to zero the coefficients without importance I will use the results to get some insights.

The evaluation is done in a 4 fold cross validation with the choice of best regularization term (C) done from default options of LogisticRegressionCV class. The metric evaluated is the ROC auc since the desbalance between the classes it's not extreme.

```
[ ]: X = df.drop("diagnosis", axis = 1)
     y = df["diagnosis"]

     X_train, X_test, y_train, y_test = train_test_split(X,y, random_state=42,
     ↪shuffle=True, stratify=y)
     cv = StratifiedKFold(n_splits = 4, shuffle=True, random_state=42)
```

```
[ ]: #Need to scale because regularization. For that, I chained the scaler and
     ↪estimator inside a Pipeline
     logit_base = LogisticRegressionCV(solver = "liblinear", cv =cv, penalty = "l1",
     ↪scoring = "roc_auc", max_iter = 1000, random_state = 42)
     baseline = make_pipeline(StandardScaler(), logit_base)
```

```
[ ]: #need to fit to get the coeficietns and C value
     baseline.fit(X_train,y_train)
```

```
[ ]: Pipeline(steps=[('standardscaler', StandardScaler()),
                      ('logisticregressioncv',
                       LogisticRegressionCV(cv=StratifiedKFold(n_splits=4,
                                                                random_state=42, shuffle=True),
                                                                max_iter=1000, penalty='l1',
                                                                random_state=42, scoring='roc_auc',
                                                                solver='liblinear'))]))
```

```
[ ]: print(f"Baseline ROC auc: {baseline[1].scores_[1][0:,4].mean()}")
      print(f"Best C: {baseline[1].C_[0]}")
```

Baseline ROC auc: 0.9942542444060662
Best C: 2.782559402207126

This is a very good cross validation score. We could suspect overfitting but as the validation was correctly done it is more likely that this high score is due to the linear separability of the classes Malign and Benign.

6 Model Selection

I also evaluated three other models (Decision Tree, Random Forest and XGboost).

For the first I use some arbitrary hyperparameters to prune the tree. As in random forest the trees are independent, a high number of estimators tends to more stable predictions (the bagging nature also contribute to that). This does not occur in boosting models like XGBoost, in which case the trees are trained in a subsequent way and each one learns from the mistakes of the previous one. In this case, a greater number of estimators tends to overfitting.

```
[ ]: #decision tree - prune the tree to avoid "infinite" growing
      dt = DecisionTreeClassifier(max_depth = 4, min_samples_leaf = 40, random_state=42)
      dt_cross_val = cross_val_score(dt, X_train, y_train, cv = cv, scoring = "roc_auc")
      print(f"Decision Tree ROC auc: {np.mean(dt_cross_val)}")
```

Decision Tree ROC auc: 0.943455970480262

```
[ ]: #random forest - high n_estimators is better
      rf = RandomForestClassifier(n_estimators = 1000 , random_state= 42)
      rf_cross_val = cross_val_score(rf, X_train, y_train, cv = cv, scoring = "roc_auc")
      print(f"Random Forest ROC auc: {np.mean(rf_cross_val)}")
```

Random Forest ROC auc: 0.9883828160143949

```
[ ]: #XGBoost - in opposite to Random Forest here more trees tend to overfit
xgb = xgb.XGBClassifier(n_estimators = 100, random_state = 42)
xgb_cross_val = cross_val_score(xgb, X_train, y_train, cv = cv, scoring =
    ↪ "roc_auc")

print(f"XGBoost ROC auc: {np.mean(xgb_cross_val)}")
```

XGBoost ROC auc: 0.9883828160143949

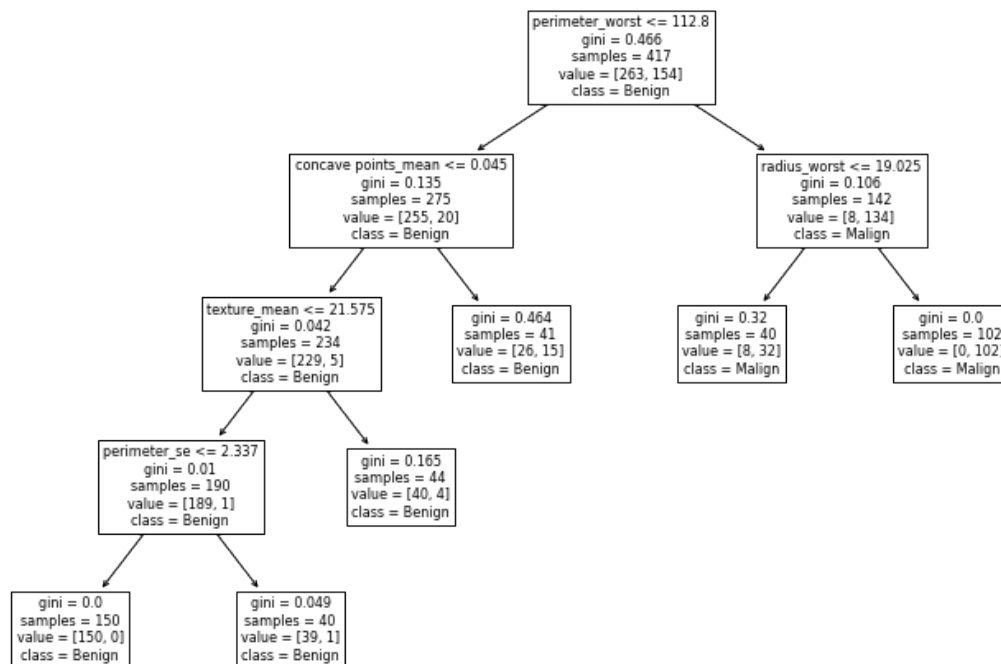
Curiously none of them (including two complex ensemble models) could beat the baseline, a simple Logistic Regression.

7 Model Interpretability

As the decision tree has been pruned, it is easy to interpret the decisions made by the model.

```
[ ]: #fitting estimator
dt.fit(X_train, y_train)

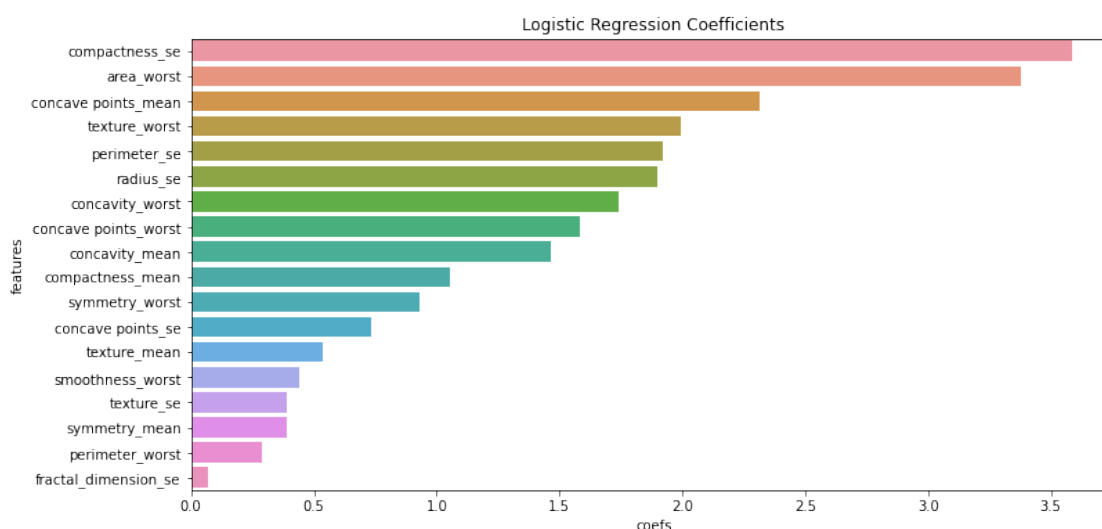
#plotting the tree
feat_dt = ["perimeter_worst", "concave points_mean", "texture_mean", "concave_
    ↪ points_worst"]
fig, ax = plt.subplots(figsize = (12,8))
plot_tree(dt, feature_names = X_train.columns, class_names =
    ↪ ["Benign", "Malign"]);
```



For example, a cell with `perimeter_worst` ≥ 112.8 and a `radius_worst` ≥ 19.025 has a 100% chance of being malignant based on 102 observations from training data.

```
[ ]: #Logistic regression Coefficients
coefs = pd.DataFrame({'features':X_train.columns,
                      'coefs': abs(baseline[1].coef_[0])})

[ ]: fig, ax = plt.subplots(figsize = (12,6))
sns.barplot(data = coefs.query('coefs > 0').sort_values(by = 'coefs', ascending_
↪ = False),
            x='coefs',
            y='features')
ax.set_title('Logistic Regression Coefficients');
```



Lasso regularization zeroed the coefficients from 9 features. Logistic Regression give too much importance to `compactness_se` and `area_worst`. We can interpret this as the following: the greater the variation in compactness between different images of the same nucleus, the greater the tendency for the cancer to be malignant. The same can be said for the area measurement.

8 Final Evaluation on Test

To evaluate on test I instantiated a new object from the class `LogisticRegression` with the `C` value defined latter on cross validationn and fitted it again on entire train set.

```
[ ]: logit = LogisticRegression(C=2.782559402207126, solver = "liblinear", penalty =_
↪ "l1", max_iter = 1000, random_state = 42)
final_model = make_pipeline(StandardScaler(), logit)
```

```
[ ]: final_model.fit(X_train, y_train)
```

```
[ ]: Pipeline(steps=[('standardscaler', StandardScaler()),  
                      ('logisticregression',  
                       LogisticRegression(C=2.782559402207126, max_iter=1000,  
                                           penalty='l1', random_state=42,  
                                           solver='liblinear'))])
```

```
[ ]: yhat = final_model.predict(X_test)
```

```
[ ]: roc_auc_score(y_test, yhat)
```


```
[ ]: 0.9501748251748252
```

We successfully managed to avoid overfitting on training since the difference in the metric was not that high

```
[ ]: print(classification_report(y_test, yhat))
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	88
1	0.96	0.92	0.94	52
accuracy			0.96	140
macro avg	0.96	0.95	0.95	140
weighted avg	0.96	0.96	0.96	140

```
[ ]: cm = confusion_matrix(y_test, yhat, labels = final_model[1].classes_)  
     disp = ConfusionMatrixDisplay(cm, display_labels = ['Benign', 'Malign'])  
     disp.plot();
```



The final result apoints to a good others metrics like f1, recall and precision for boths classes. Also, the model makes few mistakes only 4 False Negatives and 2 False Positive

9 Future Perspectives

This model can be applied to help automate the diagnosis of breast cancer as it has excellent metrics, is easily interpretable and the classes are easily distinguishable.

There are some measures that can be taken to improve it: - Treat the asymmetry in the data making them more normal. This would increase the reliability of the coefficients by improving the interpretation, although it does not affect the predictive ability. - Here I only treat outliers on the `se_` features. An alternative would be to look at the restants features. - Optimize the model to avoid false negatives by increasing the recall of class 1. This is important because this type of error is very harmful (model identifies as benign when in fact it is malignant)