

Predicting_Occurrence_of_Diabetes

December 17, 2022

1 Overview

This is my solution from the course project of IBM [Deep Learning and Reinforcement Learning](#) course. Since it's my first contact with Deep Learning, I chose to use the same dataset from the instructor (but with different analysis) which is the [Pima Indians Diabetes Database](#) retrieved from Kaggle.

Here the main objective is: **build a neural network with Keras package to predict whether or not the patients in the dataset have diabetes and compare the performance with an baseline model.**

The dataset was composed by 768 patients and 9 characteristics from them which are: - Number of times pregnant - Plasma glucose concentration a 2 hours in an oral glucose tolerance test - Diastolic blood pressure (mm Hg) - Triceps skin fold thickness (mm) - 2-Hour serum insulin (μ U/ml) - Body mass index (weight in kg/(height in m)²) - Diabetes pedigree function - Age (years) - Class variable (0 or 1)

2 Necessary packages

```
[2]: #core
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
import numpy as np

from sklearn.model_selection import train_test_split, cross_validate, \
    StratifiedKFold
from sklearn.metrics import roc_auc_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, roc_auc_score

#Auto EDA -- !pip install dataprep --
#from dataprep.eda import create_report

from keras.models import Sequential
from keras.layers import Input, Dense, Flatten, Dropout, BatchNormalization
```

```
from keras.optimizers import Adam, SGD, RMSprop
```

3 Glimpse on Data

```
[3]: PATH = "../data/"
```

```
[4]: #loading dataframe
df_raw = pd.read_csv(PATH+"diabetes.csv")

#create a copy to avoid edit raw data
df = df_raw.copy()
```

```
[5]: df.head()
```

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[ ]: df.tail()
```

```
[ ]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

```
[ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

- There are no missing value
- All features are numerical so they don't need encode

I always like to start the analysis with an auto EDA tool to give a first look on data and and get quick insights. Here I use [dataprep](#) for this purpose.

```

[ ]: #this code export a html file named diabetes_job.html with dataprep output
      ↪which can be opened locally by any browser
report = create_report(df, title='Diabetes Prediction')
report.save("diabetes_job.html")

```

Report has been saved to diabetes_job.html!

From the Auto EDA it's possible to note that: - Glucose, BloodPressure, SkinThickness, Insulin and BMI have 0 values wich is biologically impossible. - There is a desbalance between the classes, 65.1% from the patients don't have diabetes meanwhile 34.9% have.

4 Fixing Inconsistences

Since the dataset has a relatively small number of observations, drop the rows with 0 value from the collumns cited is inviable. Since the features had some skewness I impute those values with the median.

```

[ ]: inconsistent_features = ["Glucose", "BloodPressure", "SkinThickness",
      ↪"Insulin", "BMI"]

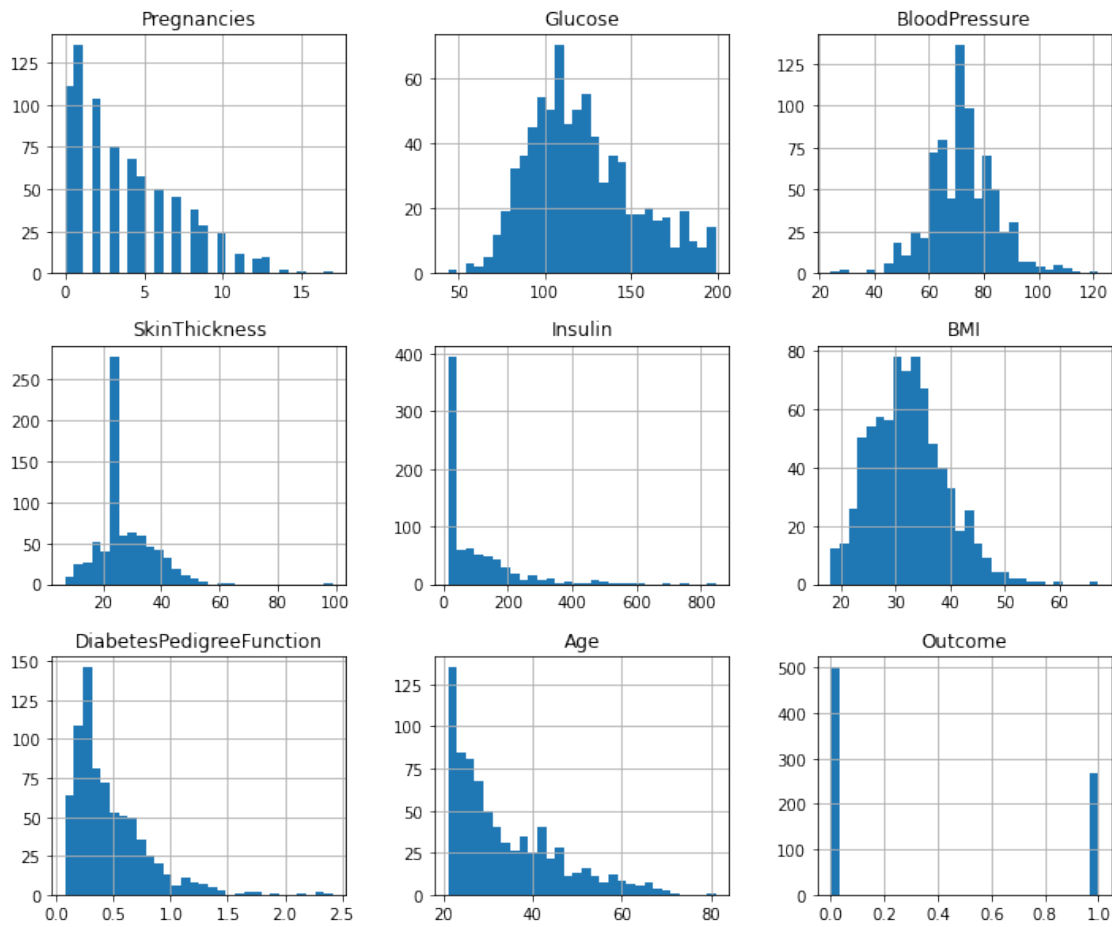
def replace_median(feature):
    df[feature] = df[feature].replace(0, df[feature].median())

for i in inconsistent_features:

```

```
replace_median(i)
```

```
[ ]: df.hist(figsize = (12,10), bins = 30);
```



```
[ ]: df.Insulin.min()
```

```
[ ]: 14.0
```

5 Machine Learning

```
[ ]: #Separating X and y features
```

```
X = df.drop("Outcome", axis = 1)
y = df["Outcome"]
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y, random_state = 42,  
→test_size = 0.3, stratify = y)
```

6 Making a Baseline

For the baseline I will train a Random Forest. Since it's a tree method don't require any type of scale or normalization.

```
[ ]: rf = RandomForestClassifier(n_estimators = 1000, random_state = 42, n_jobs = -1)

[ ]: rf.fit(X_train, y_train)

[ ]: RandomForestClassifier(n_estimators=1000, n_jobs=-1, random_state=42)

[ ]: yhat = rf.predict(X_test)

[ ]: #accuracy
accuracy_score(y_test, yhat)

[ ]: 0.7489177489177489

[ ]: #roc
roc_auc_score(y_test, yhat)

[ ]: 0.7044444444444444
```

7 Build a Single Hidden Layer Neural Network

We need to scale the data to make easier the convergence of Gradient Descendent from network

```
[ ]: scaler = StandardScaler()
X_train_scale = scaler.fit_transform(X_train)
X_test_scale = scaler.transform(X_test)

[ ]: # The input had 8 features, so the input_shape is (8,). The model had 1 hidden_
    ↪ layer, 12 hidden nodes and sigmoid activation.
    # Final layer has just one node with a sigmoid activation

model_1 = Sequential()
model_1.add(Dense(12, input_shape = (8,), activation = 'sigmoid'))
model_1.add(Dense(1, activation='sigmoid'))

model_1.compile(SGD(learning_rate = .003), "binary_crossentropy")

[ ]: yhat = (model_1.predict(X_test_scale) > 0.5).astype("int32")

8/8 [=====] - 0s 2ms/step

[ ]: #accuracy
accuracy_score(y_test, yhat)
```

```
[ ]: 0.35064935064935066
```

```
[ ]: #roc  
roc_auc_score(y_test, yhat)
```

```
[ ]: 0.5
```

The Neural Network perform very bad related the Random Forest. This should because NN is “data hungry” and we provied few data to the model.